

Computer Systems B

COMS20012

Introduction to Operating Systems and Security

OS structure

bristol.ac.uk



What is an OS?

Protection Boundary



Kernel
File System Networking
Device Drivers Processes
Virtual Memory

Hardware/Software Interface



Protection Boundaries

- Multiple privilege levels
- Different software can run with different privileges
- Processors provide at least two different modes
 - User space: how “regular” program run (e.g. steam)
 - Kernel mode: How the kernel run
- The mode determine a number of things
 - What instructions may be executed
 - How addresses are translated
 - What memory locations can be accessed (through translation)

Example Intel

- Four modes
 - Ring 0: most privileged run the kernel here
 - Ring 1/2: ignored in most environment. Can run less privileged code (e.g. third party device drivers or virtual machine monitors etc.)
 - Ring 3: where “normal” processes live
- Memory is divided in segments
 - Each segment has a privilege level (0 to 3)
 - Processor maintain a current privilege level (CPL) generally the CPL of the segment containing the instruction currently executing
 - Can read/write in segment when CPL \geq segment privilege
 - Cannot directly call code in segment where CPL $<$ segment privilege

Example MIPS

- Standard two modes processor
 - User mode: access CPU registers; flat uniform virtual memory address space
 - Kernel mode: can access memory mapping hardware and special registers

Changing Protection Level

- How do we transfer control between applications and kernel?
- When do we transfer control between applications and kernel?

How?

- Fundamental concept **trap**
- There is different types of trap which relate to the **when**

When?

- Sleeping beauty approach

- Wait for something to happens to wake up the kernel
- What might that be?
 - System calls: an application wants the operating to do something on its behalf (e.g. access some hardware)
 - Exceptions: an application does unintentionally something it should not (e.g. divide by zero, read an address it should not etc.)
 - Interrupts: An asynchronous event (e.g. I/O completion)

- Alarm clock approach

- Set a timer that generate an interrupt when it finishes

Trap

- Each type of trap is assigned a number. For example
 - System call: 1
 - Timer interrupt: 2
 - Disk interrupt: 3
- The kernel build a table that **map trap number to function** addresses to be executed
- These functions are called **trap handlers**



Example MIPS

- MIPS has 5 distinct traps, and their addresses are hardwired
 - One each for: reset, NMI (non-maskable interrupt), fast-TLB loading and debug
 - Sys161 does not support NMI and debug
 - One for everything else (software must then do further dispatch)
- Vital information are set in specific registers. For example:
 - The EPC (program counter) tell you which instruction caused an exception
 - Cause register indicate the source of the trap
 - Interrupt
 - Exception
 - System call
 - And what specific type it was
 - Status register indicate
 - The mode the processor was when it happened (i.e. user/kernel)
 - States about which kind of interrupt are enabled

Example Intel

- Hardware register
 - Traditionally called PIC (Program Interrupt Controller)
 - ... then APIC (advanced PIC)
 - ... more recently LAPIC (local APIC) one per core
- Remember
 - X86 has multiple protection level (4)
 - Cannot call code in a different level
 - This needs to be handled
- Interrupt Descriptor Table (IDT)
 - Point to special objects called gate
 - Raises the CPL
 - When returning from a gate drops CPL to its original level
- First 32 gates reserved for hardware defined traps
- Remaining available to software using the INT (interrupt) instruction.
- For system calls Linux uses a single designated INT instruction
 - Dispatch via software in the same way we saw for MIPS

Thank you

bristol.ac.uk

